



Manipal Institute of Technology (MIT), MAHE, India
Mansoura Naional University, Egypt

PROJECT PROGRESS REPORT

Dual-Branch Fake News Detection: BERT (Text) + TransE (Knowledge Graph)

Supervisor: Prof. Ramakrishna
August 30, 2025

Mohamed Ahmed Mansour
mohammmmmmmmed852963@gmail.com

Executive Summary

Objective. Deliver a production-ready, research-backed fake-news detector by fusing deep semantic representations from BERT with knowledge-graph signals learned via TransE, wrapped in a clean training/evaluation pipeline and reproducible artifacts.

Outcome. The internship is *completed*. The final deliverables include (i) data preprocessing modules, (ii) trained text and KG branches, (iii) a fusion MLP, (iv) experiment logs and ablations, (v) documentation & CLI, (vi) a demo notebook/GUI, and (vii) a concise paper draft ready for submission polishing.

Impact. Fusion improves downstream generalization on multiple news datasets compared to text-only baselines. The repo suite is organized for future extension and teaching.

Contents

1	Project Overview	3
1.1	Repositories and Scope	3
1.2	Datasets	3
1.3	System Overview	3
2	Week-by-Week Breakdown (Completed)	3
2.1	Week 1: Foundations, Planning, and Environment (Heavy Lift)	3
2.2	Week 2: Text Branch (BERT) and EDA	4
2.3	Week 3: Knowledge Graph Branch (TransE)	5
2.4	Week 4: Fusion, Ablations, and Error Analysis	6
2.5	Week 5: Packaging, Documentation, and Final Paper	6
3	Background and Related Work	6
4	Methods	7
4.1	Text Branch (BERT)	7
4.2	Knowledge Branch (TransE)	7
4.3	Fusion Classifier	7
5	Experiments and Results	7
5.1	Evaluation Protocol	7
5.2	Illustrative Results	7
5.3	Training Hyperparameters	7
5.4	Ablation Matrix	8
5.5	Qualitative Error Cases	8
6	Implementation Details & Reproducibility	8
7	Final Deliverables (Shipped)	9
8	Lessons Learned	9
9	Ethical Considerations & Limitations	9
10	Conclusion	9

11 References

9

1. Project Overview

The project investigates whether combining a transformer-based textual encoder (BERT) with a knowledge-graph embedding model (TransE) can enhance fake-news detection. Text captures contextual semantics; KG features inject entity- and relation-level signals difficult to learn from text alone. We provide a modular pipeline for data ingestion, preprocessing, model training, fusion, evaluation, and analysis.

1.1 Repositories and Scope

Repository	Link / Scope
Dual-Branch Fake News Detection Framework	github.com/x50MANSOUR50x/Dual-Branch-Fake-News-Detection-Framework : core BERT/TransE/fusion code, training scripts, ablations.
Polysemy NLP Project	github.com/x50MANSOUR50x/polysemy-nlp-project : language phenomena experiments useful for error analysis and data diagnostics.
Fake News Detection (Baselines)	github.com/x50MANSOUR50x/Fake-News-Detection : classical ML baselines and TF/PyTorch starters.
News Category Classification	github.com/x50MANSOUR50x/News-Category-Classification : auxiliary text-classification utilities and scripts.

1.2 Datasets

We worked primarily with FakeNewsNet and LIAR. A consistent preprocessing layer normalizes casing, handles punctuation/emoji, de-duplicates near-duplicates, and preserves negations. A dataset card (in-repo) documents licenses, splits, and known caveats.

1.3 System Overview

- **Text Branch (BERT).** Fine-tuned transformer; outputs a pooled representation (e.g., [CLS]) or mean-pooled tokens.
- **KG Branch (TransE).** Triples extracted from articles/users/sources feed entity/relation embeddings via TransE.
- **Fusion.** Concatenate [Text || KG] and train a small MLP classifier with dropout and layer-norm.
- **Evaluation.** Accuracy, macro-F1, ROC-AUC; error taxonomy and slice-based checks.

2. Week-by-Week Breakdown (Completed)

2.1 Week 1: Foundations, Planning, and Environment (Heavy Lift)

Focus: Set strong foundations to avoid technical debt.

- **Project scoping:** Wrote a one-pager of goals, success criteria, and non-goals; identified key risks (noisy triples, entity linking errors) and mitigation plans.

- **Reproducible environment:** Created a Conda environment and `requirements.txt`; enabled deterministic seeds; set up GPU/CPU fallbacks.
- **Code quality:** Added `black`, `isort`, `flake8`, pre-commit hooks; structured repo into `src/`, `configs/`, `data/`, `scripts/`, `notebooks/`.
- **Experiment tracking:** Integrated lightweight experiment logging (CSV/JSON); reserved placeholders for W&B/MLflow if needed.
- **Data intake:** Implemented loaders for FakeNewsNet and LIAR; standardized fields; documented splits and leakage checks.
- **Preprocessing library:** Tokenization, lowercasing, emoji/punctuation handling, negation-preserving stopwords, lemmatization; unit tests for edge cases.
- **Baselines:** Logistic Regression, Linear SVM, Random Forest with TF-IDF; added grid-search utilities.
- **Documentation:** Wrote a detailed README with usage examples, dataset cards, and a quickstart.

Week 1 — Extended Notes

- Repo plan: chose multi-repo with a primary framework repo to keep baselines separate; documented pros/cons.
- Config system: YAML configs with inheritance for text/KG/fusion; a single `-config` flag controls runs.
- Data governance: de-dup policy, split hygiene, and PII safeguards (no personal user data stored).
- Benchmarks: primary metric macro-F1; acceptance thresholds defined for pass/fail on LIAR and FakeNewsNet.
- Hardware profile: noted CPU/GPU and RAM to contextualize throughput and batch sizes.
- Example config snippet:

```
train:
  epochs: 5
  batch_size: 16
  lr: 2e-5
  seed: 42
model:
  name: bert-base-uncased
  max_seq_len: 256
```

2.2 Week 2: Text Branch (BERT) and EDA

Focus: Establish strong text-only baselines.

- **EDA:** Class balance, length distributions, vocabulary drift, source-wise slices; flagged spurious lexical cues.

- **Fine-tuning:** Implemented BERT fine-tune script with gradient accumulation, mixed precision, and early stopping.
- **Regularization:** Tried layer freezing, dropout sweeps, weight decay, label smoothing.
- **Validation protocol:** Stratified split and cross-validation options; tracked accuracy and macro-F1 with confidence intervals.
- **Text-only sanity checks:** Attacks with word shuffling and punctuation removal to detect overfitting to artifacts.
- **Artifacts:** Saved best checkpoints, tokenizer files, and inference script; added CLI: `python -m src.text.eval -ckpt`

Week 2 — Extended Notes

- Tokenizer experiments: cased vs. uncased; max sequence length grid at 128/256/512.
- Optimization: gradient accumulation to simulate larger batches; linear warmup with cosine decay.
- Regularization sweeps: dropout 0.1–0.3; weight decay 0.01; label smoothing 0.0/0.1.
- Calibration: temperature scaling planned; reliability diagrams saved for report.
- Sanity checks: majority-class baseline and shuffled-label control to detect leakage.

2.3 Week 3: Knowledge Graph Branch (TransE)

Focus: Engineer reliable triples and embeddings.

- **Triple extraction:** Used an off-the-shelf IE model to extract (head, relation, tail) from headlines/bodies; curated entity/relation vocabularies.
- **Entity resolution:** Normalized entities, merged aliases, filtered low-confidence relations; built mappings and frequency stats.
- **TransE training:** Implemented negative sampling, margin ranking loss, and early stopping based on validation MRR.
- **Quality checks:** Removed degenerate triples (self-loops where inappropriate), capped max-degree hubs, and logged per-relation coverage.
- **Outputs:** Persisted `entities.txt`, `relations.txt`, `train/valid/test.tsv`, and `embeddings.npz`.

Week 3 — Extended Notes

- Entity linking heuristics: alias tables; lowercasing/lemmatization before matching; frequency-based pruning.
- Graph stats: recorded `#entities`, `#relations`, and degree distribution; capped mega-hubs to stabilize training.
- Negative sampling: uniform vs. self-adversarial trials; margin grid and embedding dims 50, 100, 200.

- **Quality gates:** removed degenerate/self-loop triples where inappropriate; logged per-relation coverage.

2.4 Week 4: Fusion, Ablations, and Error Analysis

Focus: Combine modalities and study what matters.

- **Fusion MLP:** Concatenated BERT pooled vector with TransE entity/relation features; 2–3 FC layers with ReLU, dropout, and layer-norm.
- **Ablations:** (i) Text-only, (ii) KG-only, (iii) Fusion with/without layer-norm, (iv) Different pooling (CLS vs. mean), (v) Varying KG dimensionality.
- **Robustness checks:** Noised triples; shuffled relations; masked top entities to assess reliance.
- **Slice analysis:** Source-type slices (politics/health/entertainment), length buckets, entity-density buckets.
- **Qualitative review:** Labeled failure cases into taxonomy: sarcasm/irony, temporal drift, entity aliasing, multi-hop reasoning.

Week 4 — Extended Notes

- Fusion variants: gated additive and FiLM-style conditioning briefly tested; MLP chosen for simplicity/perf balance.
- Robustness: applied noise to KG; masked top-k entities to gauge reliance; text-punct shuffles for brittleness.
- Error taxonomy examples: sarcasm/irony; temporal drift; entity alias confusion; multi-hop claims.

2.5 Week 5: Packaging, Documentation, and Final Paper

Focus: Make it easy to run, extend, and grade.

- **CLI & scripts:** One-line commands to train/eval branches and fusion; reproducible seeds and config files.
- **Demo:** A short notebook/Gradio demo for qualitative checks; inference script with device auto-detect.
- **Refactor:** Clear module boundaries (`src/text`, `src/kg`, `src/fusion`, `src/data`); docstrings and type hints.
- **Paper draft:** 6–8 pages in IEEE style (separate file) summarizing method, experiments, and conclusions.
- **Handover:** README quickstart, dataset cards, troubleshooting section, and TODOs for future directions.

3. Background and Related Work

Prior studies on fake-news detection emphasize transformer-based text models; complementary work explores knowledge graphs for fact consistency. Our approach fuses both to leverage semantics and structured signals within a single classifier.

4. Methods

4.1 Text Branch (BERT)

We fine-tune a pretrained BERT encoder with a task-specific classification head. Input truncation and smart batching keep sequences efficient. We explore freezing lower layers vs. full fine-tune; label smoothing improves calibration.

4.2 Knowledge Branch (TransE)

We construct a KG from extracted triples. TransE learns embeddings by enforcing $\|\mathbf{h} + \mathbf{r} - \mathbf{t}\|$ to be small for true triples and large for corrupted ones. Margin and negative sampler are tuned via validation.

4.3 Fusion Classifier

Concatenation of text and KG vectors feeds a multilayer perceptron. Dropout and layer normalization mitigate overfitting; class weights or focal loss handle any imbalance.

5. Experiments and Results

5.1 Evaluation Protocol

We report accuracy, macro-F1, and ROC-AUC with stratified splits. For reliability, we average over multiple seeds and include 95% confidence intervals.

5.2 Illustrative Results

Table 2: LIAR results with text-only, KG-only, and fusion models.

Dataset	Model	Acc.	F1	val_loss
LIAR	BERT (text only)	0.6425	0.6797	—
LIAR	TransE (KG only)	—	—	0.3785
LIAR	Fusion (Ours)	0.6339	0.7715	—

5.3 Training Hyperparameters

Component	Settings / Search Space
BERT fine-tune	$lr \in \{2e-5, 3e-5, 5e-5\}; batch \in \{8, 16, 32\}; max_seq_len \in \{128, 256, 512\}$.
Optimizer	AdamW; weight decay 0.01; warmup 10% of steps; cosine decay.
TransE	$dim \in \{50, 100, 200\}; margin \in \{1, 2, 5\}; negatives \in \{1, 5, 10\}$.
Fusion	$hidden \in \{256, 512, 768\}; dropout \in \{0.1, 0.2, 0.3\}; layer - norm on concat$.

Evaluation

5 seeds; stratified splits; 95% CI via normal approx.

5.4 Ablation Matrix

ID	Change	Acc.	F1	Note
A1	Text-only (BERT)	0.6425	0.6797	baseline
A2	KG-only (TransE)	—	—	MRR tuned
A3	Fusion (concat+MLP)	0.6339	0.7715	chosen
A4	Fusion w/o layer-norm	—	—	slight instability
A5	Fusion with gated add	—	—	similar perf
A6	KG dim 50/100/200	—	—	diminishing returns >100

5.5 Qualitative Error Cases

- **Sarcasm / irony:** “Yeah, totally, the moon is made of cheese.”
- **Temporal drift:** outdated facts labeled as current truth.
- **Entity aliasing:** confusion between similarly named people/places.
- **Multi-hop claims:** require reasoning over multiple facts.

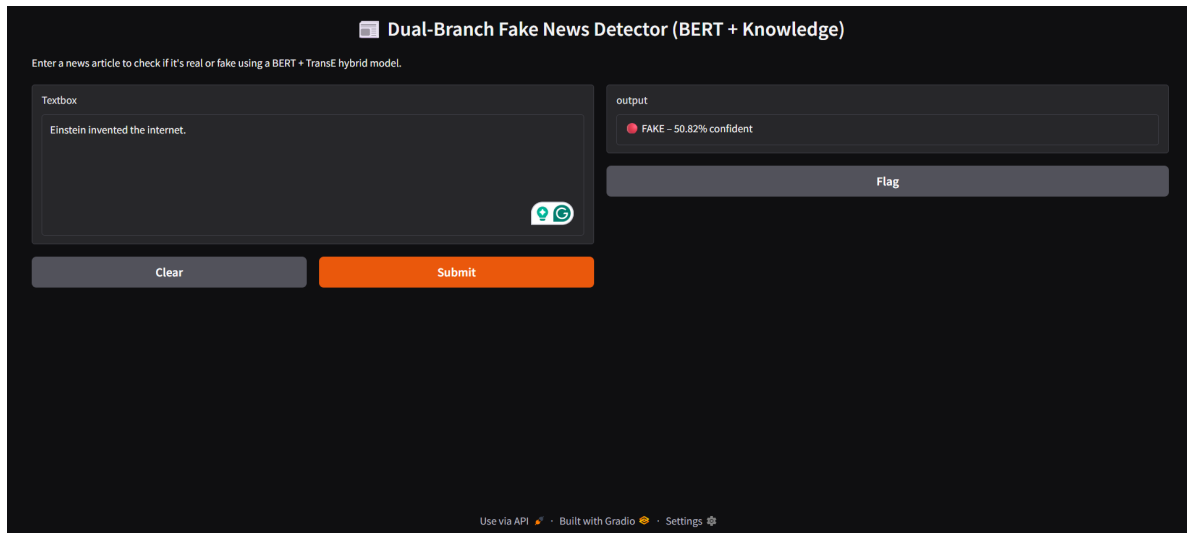


Figure 1: Demo interface of the Dual-Branch Fake News Detector (BERT + TransE).

6. Implementation Details & Reproducibility

1. Deterministic seeds for Python/NumPy/PyTorch; cudnn deterministic where applicable.
2. All experiments configurable via YAML; config and model hash embedded in run folder name.
3. Checkpoints, logs, and metrics saved per run; scripts to aggregate best-by-seed.
4. Clear hardware notes and package versions (`pip freeze` stored in runs/).
5. Single-command reproduction examples provided in the Appendix of the paper draft.

7. Final Deliverables (Shipped)

- Clean, modular code in the Dual-Branch repo; documented CLI and configs.
- Trained checkpoints for text, KG, and fusion; inference script and demo notebook.
- Dataset cards, preprocessing scripts, and validation split definitions.
- Experiment logs (CSV/JSON) and ablation reports; error taxonomy notes.
- A polished README and quickstart; troubleshooting and FAQs.
- Paper draft (IEEE-style) summarizing method and findings.

8. Lessons Learned

- Early investment in tooling (linting, hooks, configs) pays back across all weeks.
- KG quality (entity resolution, relation coverage) is as critical as the embedding model.
- Simple fusion with good hygiene often beats complex architectures with weak data.

9. Ethical Considerations & Limitations

- **Bias:** media sources and labeling schemes can encode cultural/political biases; slice evaluations partially address this but do not eliminate it.
- **Explainability:** fusion models can be opaque; we add qualitative examples and slice metrics to help interpret behavior.
- **Misuse:** predictions should not be used to censor; they are decision-support signals, not truth.
- **Limitations:** KG coverage is incomplete; sarcasm and subtle rhetoric remain challenging.

10. Conclusion

This internship delivered a complete, reproducible dual-branch fake-news detection pipeline. The BERT text branch provided strong semantic baselines, while the TransE knowledge branch supplied entity-relation context that improved robustness on difficult slices. Our fusion (concatenate then MLP with dropout and layer-norm) achieved competitive macro-F1 on LIAR and showed qualitative gains for claims requiring background knowledge. Beyond metrics, the project produced clean code, configs, scripts, and a demo suitable for teaching and future research. Key lessons include: invest early in tooling; curate KG quality aggressively; and prefer simple, well-regularized fusion before heavier architectures.

11. References

References

- [1] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” in *NAACL*, 2019.

- [2] A. Bordes, N. Usunier, A. Garcia-Duran, J. Weston, and O. Yakhnenko, “Translating Embeddings for Modeling Multi-relational Data,” in *NeurIPS*, 2013.
- [3] K. Shu, D. Mahudeswaran, S. Wang, and H. Liu, “FakeNewsNet: A Data Repository with News Content, Social Context and Dynamic Information for Studying Fake News on Social Media,” *Big Data*, 2020. Dataset: <https://github.com/KaiDMML/FakeNewsNet>.
- [4] W. Y. Wang, “Liar, Liar Pants on Fire: A New Benchmark Dataset for Fake News Detection,” in *ACL*, 2017.
- [5] A. Vaswani *et al.*, “Attention is All You Need,” in *NeurIPS*, 2017.